

```
In [1]: 1 import warnings
        2 warnings.filterwarnings('ignore')
```

```
In [2]: 1 #Importing required modules
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 import numpy as np
```

```
In [3]: 1 #Importing Churn_Modelling file
        2 churn = pd.read_csv("Churn_Modelling.csv")
```

```
In [4]: 1 #Number of rows and columns-(data points) present in the data frame
        2 churn.shape
```

Out[4]: (10000, 14)

```
In [5]: 1 #Display first 5 columns using head()
        2 churn.head()
```

Out[5]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMe
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	

```
In [6]: 1 #Display last 5 columns using tail()
        2 churn.tail()
```

Out[6]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	



```
In [7]: 1 #Checking for null values
        2 churn.isnull().sum()
```

Out[7]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

```
In [8]: 1 #List of features present in the dataset
        2 churn.columns
```

```
Out[8]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [9]: 1 #Number of features present in the dataset
        2 len(churn.columns)
```

```
Out[9]: 14
```

```
In [10]: 1 #Type of each features
         2 churn.dtypes
```

```
Out[10]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object
```

```
In [11]: 1 #Summary Statistics
        2 churn.describe()
```

Out[11]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMembe
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.00000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.51510
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.49979
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.00000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.00000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.00000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.00000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.00000

```
In [12]: 1 #Dependent variable - Exited
        2 #The Number of classes present in 'Exited'
        3 len(churn['Exited'].unique())
```

Out[12]: 2

```
In [13]: 1 #Name of two classes and the number of data points in each
        2 churn['Exited'].value_counts()
```

Out[13]: 0 7963
 1 2037
 Name: Exited, dtype: int64

Observations:

- 1.The are 1000 rows and 14 columns in the given dataset
- 2.The is no null values present.

3.Number of features present in the data frame = 14

4.The target variable is of type 'int64' which means it is numerical,: Binary flag 1 if the customer closed account with bank and 0 if the customer is retained.

5.There are 2 classes present in the target column/variable,so our problem is a classification problem

6.The dataset is an imbalanced dataset as the values of datapoints in each class is not distributed equally

Objective

Our main objective is to find whether the customer closed the account with bank (1) or the customer retained(0)

Visualizing The Data

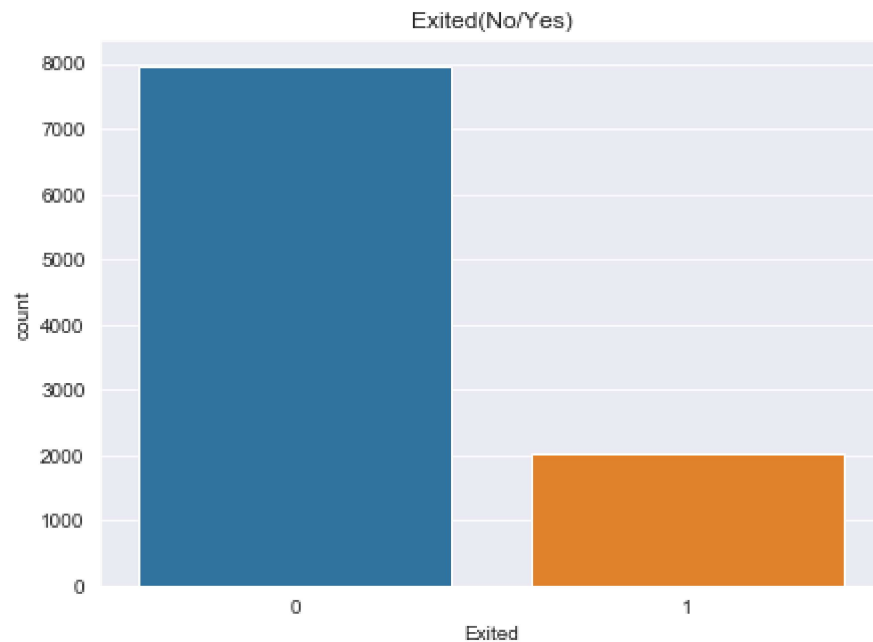
Univariate Analysis

Its all about finding out one feature which suits best with/which helps to clearly find whether the customer exited or not

Counter Plot

```
In [14]: 1 plt.figure(figsize =(7,5))
2 sns.set_style('darkgrid')
3 sns.countplot(x = 'Exited', data = churn)
4 plt.title('Exited(No/Yes)')
5 churn['Exited'].value_counts()
```

```
Out[14]: 0    7963
1     2037
Name: Exited, dtype: int64
```



The above bar graph shows that 2000 customers closed the account with the bank and about 8000 customers retained

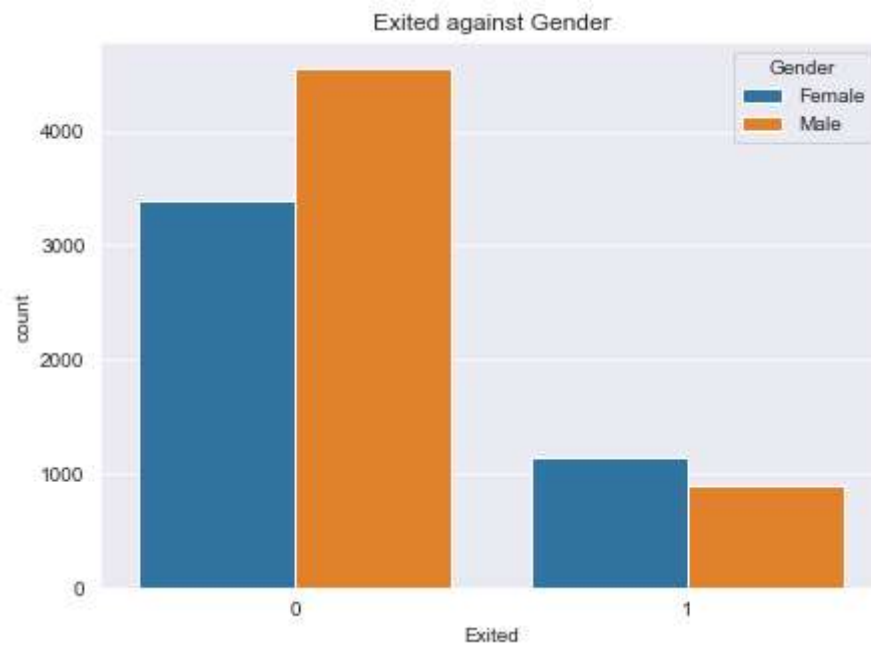
```

In [15]: 1 #Exited against Gender
          2 plt.figure(figsize =(7,5))
          3 sns.set_style('darkgrid')
          4 sns.countplot(x = 'Exited', hue = 'Gender', data = churn)
          5 plt.title('Exited against Gender');
          6 pd.DataFrame(churn.groupby(['Gender', 'Exited'])['Exited'].count())

```

Out[15]:

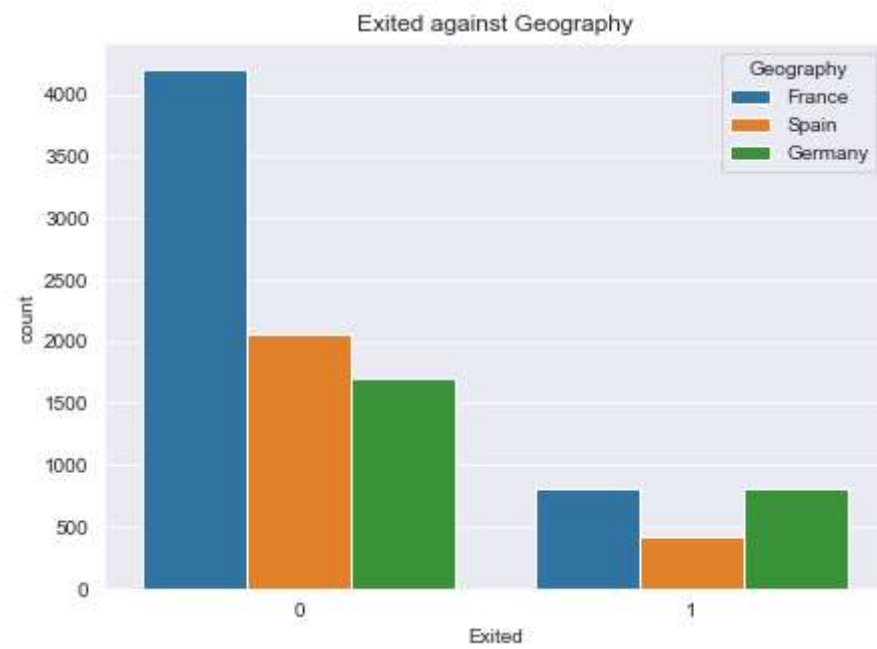
Exited		
Gender	Exited	
Female	0	3404
	1	1139
Male	0	4559
	1	898



```
In [16]: 1 #Exited against Country
2 plt.figure(figsize =(7,5))
3 sns.set_style('darkgrid')
4 sns.countplot(x = 'Exited', hue = 'Geography', data = churn)
5 plt.title('Exited against Geography');
6 pd.DataFrame(churn.groupby(['Geography', 'Exited'])['Exited'].count())
```

Out[16]:

		Exited	
Geography	Exited		
France	0	4204	
	1	810	
Germany	0	1695	
	1	814	
Spain	0	2064	
	1	413	




In []:

1

Boxplot

```
In [17]: 1 sns.boxplot(x="Exited" , y ="EstimatedSalary",data =churn)
2 plt.show()
3 sns.boxplot(x="Exited" , y ="IsActiveMember",data =churn)
4 plt.show()
5 sns.boxplot(x="Exited" , y ="HasCrCard",data =churn)
6 plt.show()
7 sns.boxplot(x="Exited" , y ="NumOfProducts",data =churn)
8 plt.show()
9 sns.boxplot(x="Exited" , y ="Balance",data =churn)
10 plt.show()
11 sns.boxplot(x="Exited" , y ="Tenure",data =churn)
12 plt.show()
13 sns.boxplot(x="Exited" , y ="Age",data =churn)
14 plt.show()
15 sns.boxplot(x="Exited" , y ="CreditScore",data =churn)
16 plt.show()
17
```



From the above box plots we can say that ,

1. Customers of age 37 to 50 are more likely to close their account with the bank


2. Customers staying with the bank for 2-3 years or 7-8 years have highest probability of closing their account with the bank.

Bivariate Analysis

It's about finding two best features which help to find the Exited status of a customer

Scatter Plot


```
In [18]: 1 sns.set_style("whitegrid");
2 sns.FacetGrid(churn, hue="Exited", size=4) \
3   .map(plt.scatter, "Age", "Exited") \
4   .add_legend();
5
6 plt.show();
7
```




Customers aged 75+ have more chance of closing their account

In [19]:

```
1 sns.set_style("whitegrid");
2 sns.FacetGrid(churn, hue="Exited", size=4) \
3   .map(plt.scatter, "Tenure", "Exited") \
4   .add_legend();
5
6 plt.show();
```



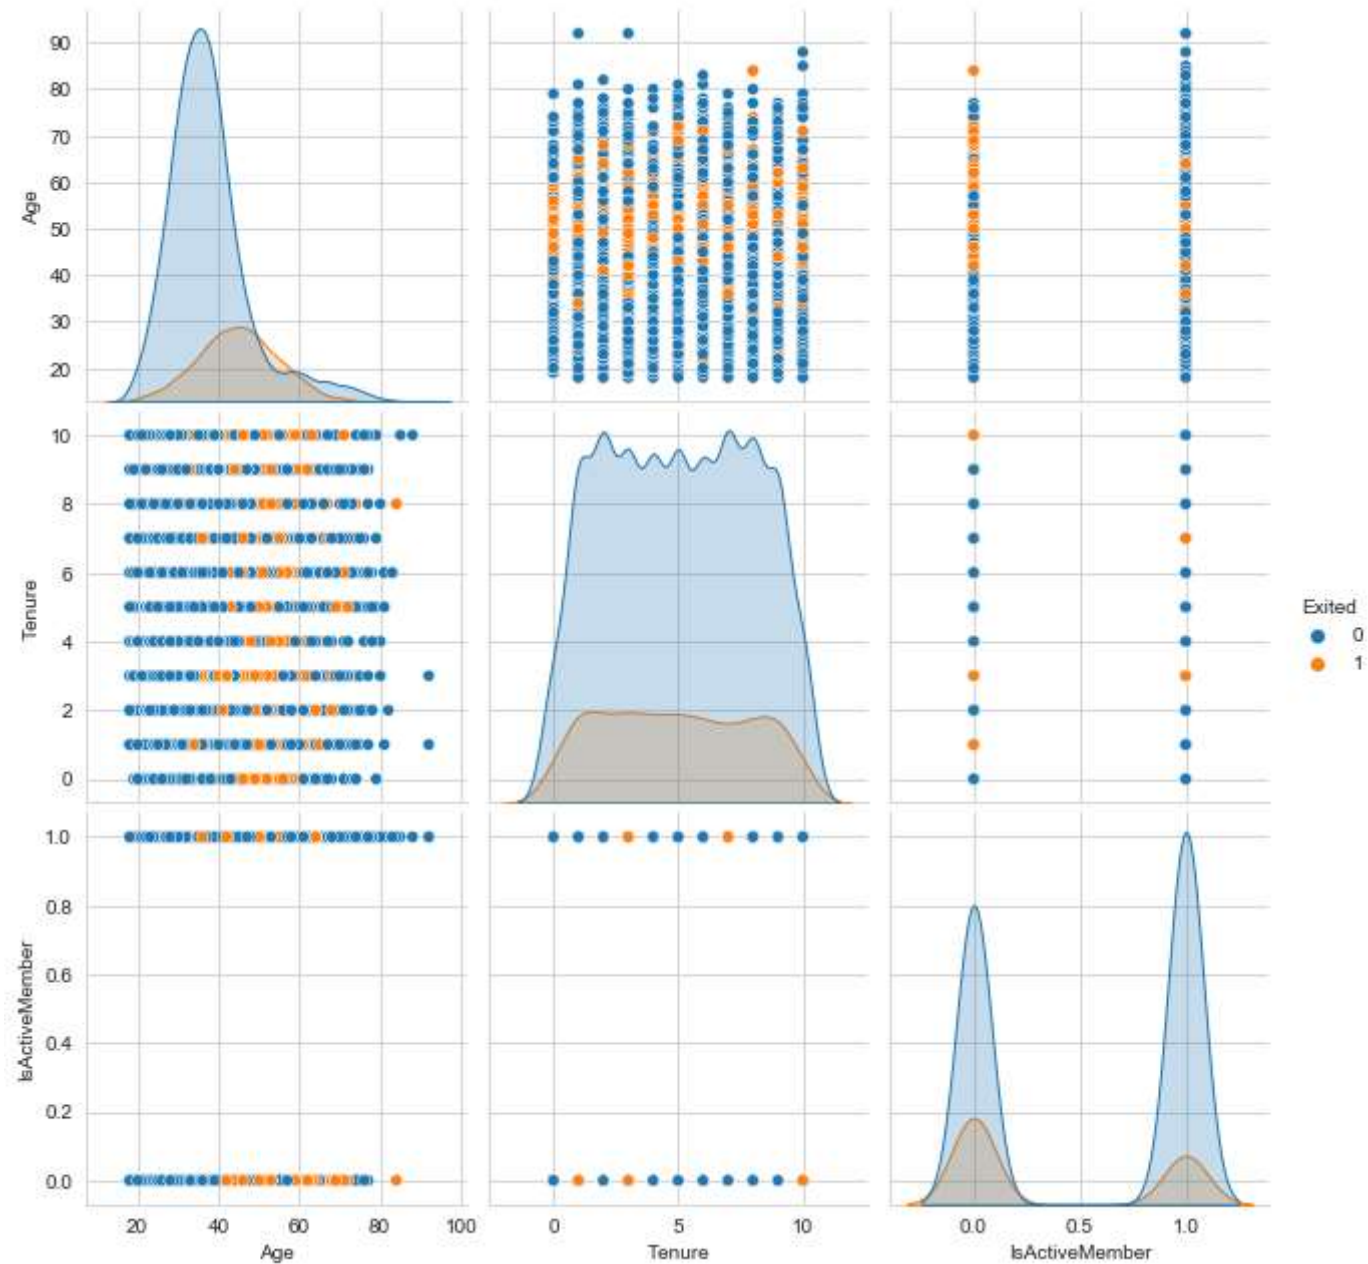
```
In [20]: 1 sns.set_style("whitegrid");
2 sns.FacetGrid(churn, hue="Exited", size=4) \
3 .map(plt.scatter, "CreditScore", "Exited") \
4 .add_legend();
5
6 plt.show();
```



Customer who has credit less than 400 are more likely to close their account with the bank

Pair Plot

```
In [21]: 1 sns.pairplot(churn, hue="Exited", size=3, vars=["Age", "Tenure", "IsActiveMember"]);
2
3 plt.show()
4
```



Descriptive Statistics

In [22]:

```
1 #Mean
2 churn.mean()
```

Out[22]:

RowNumber	5.000500e+03
CustomerId	1.569094e+07
CreditScore	6.505288e+02
Age	3.892180e+01
Tenure	5.012800e+00
Balance	7.648589e+04
NumOfProducts	1.530200e+00
HasCrCard	7.055000e-01
IsActiveMember	5.151000e-01
EstimatedSalary	1.000902e+05
Exited	2.037000e-01
dtype:	float64

```
In [23]: 1 #Median
         2 churn.median()
```

```
Out[23]: RowNumber      5.000500e+03
         CustomerId    1.569074e+07
         CreditScore    6.520000e+02
         Age           3.700000e+01
         Tenure         5.000000e+00
         Balance        9.719854e+04
         NumOfProducts  1.000000e+00
         HasCrCard      1.000000e+00
         IsActiveMember 1.000000e+00
         EstimatedSalary 1.001939e+05
         Exited         0.000000e+00
         dtype: float64
```

```
In [24]: 1 #Standard Deviation
         2 churn.std()
```

```
Out[24]: RowNumber      2886.895680
         CustomerId    71936.186123
         CreditScore    96.653299
         Age           10.487806
         Tenure         2.892174
         Balance        62397.405202
         NumOfProducts  0.581654
         HasCrCard      0.455840
         IsActiveMember 0.499797
         EstimatedSalary 57510.492818
         Exited         0.402769
         dtype: float64
```

```
In [25]: 1 #Range of Age
         2 minimum = churn['Age'].min()
         3 maximum = churn['Age'].max()
         4 range1 = maximum-minimum
         5 range1
         6
```

```
Out[25]: 74
```


In [26]:

```
1 #Quantiles
2 Q1 = churn.quantile(0.25)
3 Q2 = churn.quantile(0.50)
4 Q3 = churn.quantile(0.75)
5 print("25th Percentile :\n ",Q1)
6 print("50th Percentile :\n ",Q2)
7 print("75th Percentile :\n ",Q3)
```

25th Percentile :

RowNumber	2500.75
CustomerId	15628528.25
CreditScore	584.00
Age	32.00
Tenure	3.00
Balance	0.00
NumOfProducts	1.00
HasCrCard	0.00
IsActiveMember	0.00
EstimatedSalary	51002.11
Exited	0.00

Name: 0.25, dtype: float64

50th Percentile :

RowNumber	5.000500e+03
CustomerId	1.569074e+07
CreditScore	6.520000e+02
Age	3.700000e+01
Tenure	5.000000e+00
Balance	9.719854e+04
NumOfProducts	1.000000e+00
HasCrCard	1.000000e+00
IsActiveMember	1.000000e+00
EstimatedSalary	1.001939e+05
Exited	0.000000e+00

Name: 0.5, dtype: float64


75th Percentile :

RowNumber	7.500250e+03
CustomerId	1.575323e+07
CreditScore	7.180000e+02
Age	4.400000e+01
Tenure	7.000000e+00
Balance	1.276442e+05
NumOfProducts	2.000000e+00

```
HasCrCard      1.000000e+00
IsActiveMember 1.000000e+00
EstimatedSalary 1.493882e+05
Exited         0.000000e+00
Name: 0.75, dtype: float64
```

Outliers

```
In [27]: 1 #Age against Exited
        2
        3 sns.boxplot(x="Exited" , y ="Age",data =churn)
        4 plt.show()
```



The points above the upper and lower whiskers are called a outliers


We can use IQR to remove these outliers

```
In [28]: 1 Q1 = churn['Age'].quantile(0.25)
2 Q2 = churn['Age'].quantile(0.50)
3 Q3 = churn['Age'].quantile(0.75)
4 IQR = Q3 - Q1
5 print("25th Percentile :\n ",Q1)
6 print("75th Percentile :\n ",Q3)
7 print("IQR: ",IQR)
```

```
25th Percentile :
32.0
75th Percentile :
44.0
IQR: 12.0
```

```
In [29]: 1 #Removing Outliers
2 churn1 = churn[(churn['Age']>=Q1-(1.5*IQR)) & (churn['Age']<= Q3+(1.5*IQR))]
```

```
In [30]: 1 #Age after removing outliers
2 sns.boxplot(x="Exited" , y ="Age",data =churn1)
3 plt.show()
```



Encoding of Categorical Variables

In [31]:

```
1 churn2 = churn
```

In [32]:

```
1 #Encoding Age  
2 #Female-1 and Male-0  
3 for i in churn2.index:  
4     if(churn2['Gender'][i]=="Female"):  
5         churn2['Gender'][i]=1  
6     else:  
7         churn2['Gender'][i]=0
```

In [33]:

```
1 #After Encoding  
2 churn2['Gender'][0:40]
```

Out[33]:

```
0      1  
1      1  
2      1  
3      1  
4      1  
5      0  
6      0  
7      1  
8      0  
9      0  
10     0  
11     0  
12     1  
13     1  
14     1  
15     0  
16     0  
17     1  
18     0  
19     1  
20     0  
21     1  
22     1  
23     0  
24     1  
25     0  
26     0  
27     0  
28     1  
29     0  
30     1  
31     0  
32     0  
33     1  
34     1  
35     1  
36     0  
37     0  
38     0
```

```
39    0
Name: Gender, dtype: object
```

```
In [34]: 1 churn2['Geography'].value_counts()
```

```
Out[34]: France      5014
Germany    2509
Spain      2477
Name: Geography, dtype: int64
```

```
In [35]: 1 #Encoding Geography
2 #France - 1
3 #Germany - 2
4 #Spain - 3
5 for i in churn2['Geography'].index:
6     if(churn2['Geography'][i]=="France"):
7         churn2['Geography'][i]=1
8     elif(churn2['Geography'][i]=="Germany"):
9         churn2['Geography'][i]=2
10    else:
11        churn2['Geography'][i]=3
```

```
In [36]: 1 #After Encoding 'Geography'
2 churn2['Geography']
```

```
Out[36]: 0      1
1      3
2      1
3      1
4      3
..
9995   1
9996   1
9997   1
9998   2
9999   1
Name: Geography, Length: 10000, dtype: object
```

In [37]:

```
1 churn2
```

Out[37]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMemb
1	15634602	Hargrave	619	1	1	42	2	0.00	1	1	
2	15647311	Hill	608	3	1	41	1	83807.86	1	0	
3	15619304	Onio	502	1	1	42	8	159660.80	3	1	
4	15701354	Boni	699	1	1	39	1	0.00	2	0	
5	15737888	Mitchell	850	3	1	43	2	125510.82	1	1	
...
9996	15606229	Obijiaku	771	1	0	39	5	0.00	2	1	
9997	15569892	Johnstone	516	1	0	35	10	57369.61	1	1	
9998	15584532	Liu	709	1	1	36	7	0.00	1	0	
9999	15682355	Sabbatini	772	2	0	42	3	75075.31	2	1	
10000	15628319	Walker	792	1	1	28	4	130142.79	1	1	

10 rows × 12 columns



Splitting Dependent and Independent variables

In [38]:

```
1 #Deleting unnecessary features
2 churn2.drop(['RowNumber', 'CustomerId', 'Surname'],axis = 1,inplace = True)
```

In [39]:

```
1 churn2
```

Out[39]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	1	1	42	2	0.00	1	1	1	101348.88	1
1	608	3	1	41	1	83807.86	1	0	1	112542.58	0
2	502	1	1	42	8	159660.80	3	1	0	113931.57	1
3	699	1	1	39	1	0.00	2	0	0	93826.63	0
4	850	3	1	43	2	125510.82	1	1	1	79084.10	0
...
9995	771	1	0	39	5	0.00	2	1	0	96270.64	0
9996	516	1	0	35	10	57369.61	1	1	1	101699.77	0
9997	709	1	1	36	7	0.00	1	0	1	42085.58	1
9998	772	2	0	42	3	75075.31	2	1	0	92888.52	1
9999	792	1	1	28	4	130142.79	1	1	0	38190.78	0

10000 rows × 11 columns

In [40]:

```
1 #Dependent variable
2 dependent = churn2.iloc[0:10000, 10:11]
```



```
In [41]: 1 dependent
```

```
Out[41]:
```

Exited	
0	1
1	0
2	1
3	0
4	0
...	...
9995	0
9996	0
9997	1
9998	1
9999	0

10000 rows × 1 columns

```
In [42]: 1 #Independent Variable
          2 independent = churn2.iloc[0:10000, 0:10]
```

```
In [43]: 1 independent
```

Out[43]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	1	1	42	2	0.00	1	1	1	101348.88
1	608	3	1	41	1	83807.86	1	0	1	112542.58
2	502	1	1	42	8	159660.80	3	1	0	113931.57
3	699	1	1	39	1	0.00	2	0	0	93826.63
4	850	3	1	43	2	125510.82	1	1	1	79084.10
...
9995	771	1	0	39	5	0.00	2	1	0	96270.64
9996	516	1	0	35	10	57369.61	1	1	1	101699.77
9997	709	1	1	36	7	0.00	1	0	1	42085.58
9998	772	2	0	42	3	75075.31	2	1	0	92888.52
9999	792	1	1	28	4	130142.79	1	1	0	38190.78

10000 rows × 10 columns

Splitting of Train and Test

```
In [44]: 1 from sklearn.model_selection import train_test_split
```

```
In [45]: 1 x_train,x_test,y_train,y_test = train_test_split(independent,dependent,test_size = 0.2 ,random_state = 0)
```

```
In [46]: 1 x_train.shape
```

Out[46]: (8000, 10)

```
In [47]: 1 x_test.shape
```

```
Out[47]: (2000, 10)
```

```
In [48]: 1 y_train.shape
```

```
Out[48]: (8000, 1)
```

```
In [49]: 1 y_test.shape
```

```
Out[49]: (2000, 1)
```

```
In [50]: 1 x_train.head()
```

```
Out[50]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
7389	667	3	1	34	5	0.00	2	1	0	163830.64
9275	427	2	0	42	1	75681.52	1	1	1	57098.00
2995	535	1	1	29	2	112367.34	1	1	0	185630.76
5316	654	3	0	40	5	105683.63	1	1	0	173617.09
356	850	3	1	57	8	126776.30	2	1	1	132298.49

Scaling the Independent Variable - Feature Scaling

```
In [51]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [52]: 1 sc = StandardScaler()
```

```
In [53]: 1 x_train = sc.fit_transform(x_train)
```

```
In [54]: 1 print(x_train)
```

```
[[ 0.16958176  1.51919821  1.09168714 ...  0.64259497 -1.03227043
   1.10643166]
 [-2.30455945  0.3131264  -0.91601335 ...  0.64259497  0.9687384
  -0.74866447]
 [-1.19119591 -0.89294542  1.09168714 ...  0.64259497 -1.03227043
   1.48533467]
 ...
 [ 0.9015152  -0.89294542 -0.91601335 ...  0.64259497 -1.03227043
   1.41231994]
 [-0.62420521  1.51919821  1.09168714 ...  0.64259497  0.9687384
   0.84432121]
 [-0.28401079  0.3131264  1.09168714 ...  0.64259497 -1.03227043
   0.32472465]]
```

```
In [55]: 1 x_test = sc.transform(x_test)
```

```
In [56]: 1 x_test[0:5,:]
```

```
Out[56]: array([[ -0.55204276,  0.3131264 ,  1.09168714, -0.36890377,  1.04473698,
                   0.8793029 , -0.92159124,  0.64259497,  0.9687384 ,  1.61085707],
                 [-1.31490297, -0.89294542,  1.09168714,  0.10961719, -1.031415 ,
                   0.42972196, -0.92159124,  0.64259497, -1.03227043,  0.49587037],
                 [ 0.57162971,  1.51919821,  1.09168714,  0.30102557,  1.04473698,
                   0.30858264, -0.92159124,  0.64259497,  0.9687384 , -0.42478674],
                 [ 1.41696129, -0.89294542, -0.91601335, -0.65601634, -0.33936434,
                   0.57533623, -0.92159124, -1.55619021, -1.03227043, -0.18777657],
                 [ 0.57162971,  0.3131264 , -0.91601335, -0.08179119,  0.00666099,
                   1.38961097,  0.8095029 ,  0.64259497,  0.9687384 ,  0.61684179]])
```